



AI Email Summary Agent

Complete Documentation & Implementation Guide

Make.com Automation

Gmail API

HTTP Module

AI-Powered

Version 1.0 | May 2026 | Confidential

Table of Contents

1.	Introduction	3
	1.1 Project Overview	3
	1.2 Problem Statement	3
	1.3 Objectives	4
2.	System Architecture	4
	2.1 High-Level Workflow	4
	2.2 Technology Stack	5
	2.3 Data Flow Diagram	5
3.	Make.com Platform Overview	6
	3.1 What is Make.com?	6
	3.2 Scenarios & Modules	6
	3.3 Why Make.com for Email Automation?	7
4.	Gmail Integration	7
	4.1 Gmail Module Setup	7
	4.2 Watch Emails Trigger	8
	4.3 Fetch Email Content	8
	4.4 Gmail Permissions & Scopes	9
5.	HTTP Module Configuration	9
	5.1 Understanding the HTTP Module	9
	5.2 API Request Setup	10
	5.3 Request Headers & Body	10
	5.4 Handling API Responses	11
6.	AI Summarization Engine	11
	6.1 AI Model Selection	11
	6.2 Prompt Engineering	12
	6.3 Summary Output Formats	13
7.	End-to-End Workflow Walkthrough	13

	7.1 Trigger: New Email Received	13
	7.2 Email Parsing	14
	7.3 HTTP Call to AI API	14
	7.4 Summary Delivery	14
8.	Error Handling & Reliability	15
9.	Security & Privacy	16
10.	Testing & Validation	17
11.	Performance & Optimization	18
12.	Troubleshooting Guide	19
13.	Advanced Enhancements	20
14.	Use Cases & Business Impact	21
15.	Conclusion	22

1. Introduction

1.1 Project Overview

The **AI Email Summary Agent** is an intelligent automation solution built on **Make.com** (formerly Integromat) that automatically monitors a Gmail inbox, extracts email content, sends it to an AI language model via an HTTP API call, and returns a concise, human-readable summary to the user. This eliminates the need to manually read through lengthy emails, newsletters, or support threads, saving hours of time every week.

The agent runs continuously in the cloud — no servers to manage, no code to deploy — and integrates directly with Gmail using Make.com's native Gmail modules alongside the flexible HTTP module for AI API calls.

1.2 Problem Statement

Did you know? The average professional spends **28% of their workday** reading and answering emails — roughly 2.6 hours every single day. For executives, this figure climbs even higher. AI-powered email summarization directly addresses this productivity drain.

Modern inboxes are overwhelmed. Marketing newsletters, long discussion threads, forwarded reports, and automated notifications all compete for attention. Key action items are buried inside multi-paragraph emails. Decision-makers need the essence — not the entire email — to act quickly and confidently.

1.3 Objectives

- **Automate email monitoring** using Gmail Watch triggers in Make.com.
- **Extract structured data** — sender, subject, body, timestamp — from each new email.
- **Integrate AI summarization** via HTTP POST requests to an LLM API endpoint.
- **Deliver summaries** back to the user through Gmail reply, label, or separate notification.
- **Handle errors gracefully** with fallback routes and retry logic inside Make.com.
- **Maintain security** by using encrypted API keys and OAuth 2.0 for Gmail access.

2. System Architecture

2.1 High-Level Workflow

The agent follows a five-stage pipeline. Each stage corresponds to one or more Make.com modules configured inside a single Scenario:

Stage 1	Trigger	Gmail Watch Emails module polls the inbox every 1–15 minutes for new messages matching a filter.
---------	---------	--

Stage 2	Parse	The email subject, sender, HTML/plain-text body, and metadata are extracted and cleaned.
Stage 3	Summarize	An HTTP module sends the cleaned text to an AI API (OpenAI, Claude, Gemini, etc.) with a crafted prompt.
Stage 4	Process	The AI response is parsed; the summary text is extracted from the JSON payload.
Stage 5	Deliver	The summary is sent back — as a Gmail reply, a label, a Slack message, or saved to Google Drive.

2.2 Technology Stack

Component	Technology	Role
Automation Platform	Make.com	Scenario orchestration, scheduling, error handling
Email Provider	Gmail (Google Workspace)	Email source, OAuth 2.0 authentication
HTTP Client	Make.com HTTP Module	REST API calls to AI endpoint
AI Engine	OpenAI GPT / Anthropic Claude / Gemini	Natural language summarization
Auth	OAuth 2.0 + API Keys	Secure credential management
Data Format	JSON / HTML / Plain Text	Email body and API payloads

2.3 Data Flow Diagram

■ **Gmail Inbox** → [Watch Emails Module] → Email Object (subject, from, body, date) → [HTTP Module: POST /v1/messages] → AI API → JSON Response → [Parse Module] → Summary Text → [Gmail Send Module] → ■ **User Notification**

3. Make.com Platform Overview

3.1 What is Make.com?

Make.com is a visual, no-code/low-code integration platform that lets you connect apps and automate workflows through a drag-and-drop interface. Workflows — called **Scenarios** — are composed of **Modules**: pre-built connectors for hundreds of services including Gmail, Google Sheets, Slack, HTTP calls, JSON parsing, data manipulation, and much more.

Make.com executes Scenarios on a schedule (every minute, hourly, on-demand) or in response to real-time webhooks. Its visual debugger shows exactly what data flows between modules, making troubleshooting intuitive even for non-developers.

3.2 Scenarios & Modules

A **Scenario** is a visual automation chain. Each circle in the Make.com canvas represents one **Module**. Modules are connected left-to-right; data output from one module becomes the input for the next. Our AI Email Summary Agent uses the following modules in sequence:

#	Module Name	Type	Purpose
1	Gmail – Watch Emails	Trigger	Detects new emails in specified label/folder
2	Gmail – Get an Email	Action	Retrieves full email content including HTML body
3	Text Parser / Set Variable	Tool	Strips HTML tags, truncates body to token limit
4	HTTP – Make a Request	HTTP	POSTs cleaned body to the AI REST API
5	JSON – Parse JSON	Tool	Extracts summary field from AI API response
6	Gmail – Send an Email	Action	Sends the AI summary back to the user

3.3 Why Make.com for Email Automation?

- **No server maintenance** — Make.com runs entirely in the cloud. Zero DevOps required.
- **Native Gmail connector** — OAuth 2.0 authentication is handled by Make.com; no manual token refresh.
- **Visual debugging** — Every execution shows a step-by-step log with input/output bundles.
- **Error handling** — Built-in retry, ignore, rollback, and break directives per module.
- **Scalable operations** — From 1,000 to 10,000+ emails per month, just upgrade the plan.
- **HTTP flexibility** — The HTTP module works with any REST API, giving AI model freedom.

4. Gmail Integration

4.1 Gmail Module Setup

To add the Gmail module in Make.com, click the + button in your Scenario canvas, search for **Gmail**, and select the appropriate action. You will be prompted to create a Gmail connection by signing in with Google and granting the required permissions. Make.com stores the OAuth tokens securely in its credential vault.

Tip: Use a dedicated Gmail account or a Google Workspace service account for the agent rather than your personal inbox. This isolates the automation and avoids accidental access to personal emails.

4.2 Watch Emails Trigger

The **Gmail – Watch Emails** module is the entry point of the Scenario. It polls Gmail at a configurable interval and fires once for each new email that matches your criteria.

Polling Interval	1 min (Pro) / 5 min (Free) / Custom with webhooks
Label Filter	INBOX, or any custom Gmail label (e.g., 'To Summarize')
Search Query	Gmail search syntax, e.g. is:unread from:newsletter@company.com

Max Results	1–100 emails per execution cycle (default: 1 for real-time)
Mark as Read	Optional — auto-mark processed emails to avoid re-triggering
From Filter	Restrict to specific senders or domains for targeted summarization

4.3 Fetch Email Content

After the trigger fires, use the **Gmail – Get an Email** module with the Message ID from the trigger bundle. This retrieves the complete email including:

- **Subject** — The email subject line, used in the AI prompt context.
- **From** — Sender name and email address.
- **To / CC / BCC** — Recipients.
- **Date** — ISO timestamp of receipt.
- **Body (HTML)** — Full HTML-formatted email body.
- **Body (Plain Text)** — Plain text version, preferred for AI input.
- **Attachments** — List of attachments (names, sizes, MIME types).
- **Thread ID** — For replying in the same email thread.

Best Practice: Always prefer **Body (Plain Text)** over HTML when feeding content to the AI API. HTML contains tags, inline styles, and tracking pixels that consume tokens without contributing to meaning. If plain text is empty, use a Text Parser module to strip HTML tags from the HTML body before passing it to the HTTP module.

4.4 Gmail Permissions & Scopes

Make.com requests the following Gmail OAuth scopes on your behalf. These are the minimum required scopes for the agent to function:

OAuth Scope	Permission	Required For
gmail.readonly	Read all email messages and settings	Watch Emails, Get Email
gmail.send	Send emails on your behalf	Send summary email reply
gmail.modify	Apply labels, mark read/unread	Mark as read, add label
gmail.labels	Create and manage labels	Auto-label summarized emails

5. HTTP Module Configuration

5.1 Understanding the HTTP Module

The **HTTP – Make a Request** module is the most powerful generic module in Make.com. It allows you to make any HTTP/HTTPS request — GET, POST, PUT, PATCH, DELETE — to any REST API endpoint. For the AI Email Summary Agent, this module handles the communication between Make.com and the AI language model API.

Unlike dedicated app connectors, the HTTP module gives you full control over request headers, body format, authentication, SSL settings, and response parsing — making it compatible with **any AI provider**: OpenAI, Anthropic Claude, Google Gemini, Mistral, Cohere, or a self-hosted model.

5.2 API Request Setup

Method	POST
URL	https://api.openai.com/v1/chat/completions (or your chosen AI provider)
Body Type	Raw / JSON
Content-Type	application/json
Authorization	Bearer {YOUR_API_KEY} — stored in Make.com's encrypted key vault
Timeout	30–60 seconds recommended for AI inference
Parse Response	Yes — enables Make.com to auto-map JSON fields from the response

5.3 Request Headers & Body

Configure the following headers in the HTTP module:

Header Key	Header Value	Notes
Content-Type	application/json	Required for JSON body
Authorization	Bearer {{apiKey}}	Replace with Make.com variable
Accept	application/json	Ensures JSON response format
User-Agent	MakeEmailAgent/1.0	Optional, for API logging

Example JSON request body sent by the HTTP module to OpenAI:


```
{
  "model": "gpt-4o-mini",
  "max tokens": 300,
  "temperature": 0.3,
  "messages": [
    {
      "role": "system",
      "content": "You are an expert email summarizer. Be concise."
    },
    {
      "role": "user",
      "content": "Summarize this email:\n\n{{emailBody}}"
    }
  ]
}
```

5.4 Handling API Responses

When **Parse Response** is enabled, Make.com automatically maps the AI API's JSON response into mappable variables. For OpenAI, the summary text is located at:

```
// OpenAI response path:
choices[0].message.content

// Anthropic Claude response path:
content[0].text

// Google Gemini response path:
candidates[0].content.parts[0].text
```

Use a **JSON – Parse JSON** module or Make.com's built-in map function to extract the summary string. Then pass this string to the Gmail Send module.

6. AI Summarization Engine

6.1 AI Model Selection

The choice of AI model impacts summary quality, speed, and cost. Below is a comparison of popular models suitable for email summarization:

Model	Provider	Context Window	Best For	Cost (approx.)
GPT-4o Mini	OpenAI	128K tokens	Best balance of quality/cost	\$0.15/1M tokens in
GPT-4o	OpenAI	128K tokens	High accuracy, complex emails	\$2.50/1M tokens in

Claude Haiku 3	Anthropic	200K tokens	Ultra-fast, very cheap	\$0.25/1M tokens in
Claude Sonnet 3.5	Anthropic	200K tokens	Excellent quality summaries	\$3.00/1M tokens in
Gemini 1.5 Flash	Google	1M tokens	Very long email threads	\$0.075/1M tokens in
Mistral Small	Mistral	32K tokens	EU data residency option	\$0.20/1M tokens in

Recommendation: Start with **GPT-4o Mini** or **Claude Haiku 3** for cost-effective summarization. Upgrade to **GPT-4o** or **Claude Sonnet** for high-stakes emails requiring nuanced understanding. A typical email body is 200–500 tokens — cost per summary is fractions of a cent.

6.2 Prompt Engineering

The quality of AI summaries depends heavily on the **system prompt** and **user prompt** design. Below are three proven prompt templates for different use cases:

Template A — General Summary Prompt

```
SYSTEM:
You are a professional email assistant. Summarize emails clearly and concisely.
Output format:
• Subject: [one line]
• From: [sender name and role if identifiable]
• Key Points: [3-5 bullet points]
• Action Required: [Yes/No + what action, or None]
• Tone: [Professional / Casual / Urgent / Promotional]

USER:
Please summarize the following email:
---
{{emailBody}}
---
```

Template B — Action-Focused Prompt

Template C — Short Tweet-Style Summary

```
SYSTEM:
Summarize the email in 1-2 sentences max. Be direct. No fluff.

USER:
{{emailBody}}
```

6.3 Summary Output Formats

Depending on how you plan to use the summaries, configure the AI to respond in different formats:

- **Plain Text** — Simple paragraph summary, easy to email or display in notifications.
- **Bullet Points** — Structured list of key points, ideal for quick scanning.
- **JSON Object** — Machine-readable format, useful for storing in Google Sheets or a database.
- **HTML Formatted** — Rich text with bold headings, suitable for email replies.

```
SYSTEM:
Extract only actionable information from the email.
Respond in JSON format: {"summary": "...", "action": "...", "deadline": "...", "priority": "high/medium/low"}

USER:
Email subject: {{subject}}
Email from: {{from}}
Email body: {{body}}
```

- **Markdown** — Lightweight markup, useful for Notion or Slack integration.

7. End-to-End Workflow Walkthrough

7.1 Trigger: New Email Received

The Scenario begins the moment a new email arrives in the monitored Gmail label. Make.com's Gmail Watch Emails module checks the mailbox at the configured interval. When a new message is found:

- Make.com creates a new **Bundle** (a data package) containing the email metadata.
- The Bundle includes: messageId, threadId, subject, from, date.
- This Bundle is passed downstream to the next module.
- If multiple emails arrive simultaneously, each is processed as a **separate Bundle**.

7.2 Email Parsing

The **Gmail – Get an Email** module receives the messageId and fetches the complete message from Gmail's API. The email body is then cleaned:

Step 1 – Get Email	Use 'Gmail: Get an Email' module with messageId from trigger bundle.
Step 2 – Extract Body	Map the 'Body (Plain Text)' field. If empty, map 'Body (HTML)'.

Step 3 – Strip HTML	Use 'Text Parser: Replace' or a regex to remove patterns if HTML.
Step 4 – Truncate	Limit body to 3000–4000 characters to stay within AI token limits.
Step 5 – Set Variables	Store subject, from, cleaned body in named variables for the HTTP call.

7.3 HTTP Call to AI API

The HTTP module fires a POST request with the cleaned email content embedded in the prompt. The complete request cycle takes 2–8 seconds depending on email length and AI provider response time. Make.com waits for the response before proceeding.

Token Tip: A typical email of 300 words is approximately 400 tokens. With a 300-token max_tokens limit on the response, the total API cost per email is under \$0.001 (less than 0.1 cents) using GPT-4o Mini — making this extremely cost-effective even at high volumes.

7.4 Summary Delivery

After extracting the summary text from the API response, Make.com delivers it through one or more of the following configured output modules:

Delivery Method	Module Used	Configuration
Email Reply	Gmail – Send an Email	Reply to original thread with AI summary as body
Label + Archive	Gmail – Modify Labels	Add 'Summarized' label, remove INBOX label
Google Sheet Log	Google Sheets – Add a Row	Log sender, subject, summary, timestamp
Slack Notification	Slack – Create a Message	Post summary to #email-digest channel
Google Doc	Google Docs – Create a Doc	Save full summary as a dated Google Doc

8. Error Handling & Reliability

Production automation requires robust error handling. Make.com provides four error directives that can be applied individually to any module in the Scenario:

Directive	Behavior	Use Case for Email Agent
Ignore	Skip the failed module and continue	If AI API is down, skip summary and deliver raw subject
Retry	Retry the module N times with delay	Transient API rate limits (429 errors) — retry 3x
Break	Stop scenario, mark execution as Incomplete	Gmail auth failure — needs manual intervention
Rollback	Undo all changes since last commit point	Not typically needed for email agents

8.1 Common Errors and Fixes

■ Gmail 401 Unauthorized

Cause	OAuth token expired or revoked
Fix	Re-authenticate Gmail connection in Make.com settings

■ HTTP 429 Too Many Requests

Cause	AI API rate limit exceeded
Fix	Add Retry directive with exponential backoff (10s, 30s, 60s)

■ HTTP 500 Server Error

Cause	AI API internal error
Fix	Use Ignore directive + send fallback notification email

■ Empty Email Body

Cause	Email is image-only or attachment-only
Fix	Add a filter: 'Continue if Body (Plain Text) is not empty'

■ Token Limit Exceeded

Cause	Email body too long for AI context
Fix	Truncate body to 3000 chars before HTTP call

■ JSON Parse Error

Cause	AI returned malformed response
Fix	Wrap parse in a Try/Catch using an error handler route

8.2 Scenario Execution Monitoring

Make.com's **Execution History** panel shows every run with status (Success / Warning / Error / Incomplete), processing time, number of operations consumed, and full input/output data for each module. Set up **Email Notifications** in Scenario settings to be alerted immediately when a Scenario fails.

9. Security & Privacy

Email data is sensitive. The AI Email Summary Agent is designed with security-first principles across all integration points:

9.1 Credential Management

- **OAuth 2.0 for Gmail** — Make.com handles token storage and refresh. API keys are never exposed in Scenario logic.
- **Encrypted Key Vault** — Store AI API keys in Make.com's built-in encrypted data store, not in hardcoded variables.
- **Least Privilege** — Request only the Gmail OAuth scopes actually needed (e.g., don't request gmail.compose if you only read).
- **Key Rotation** — Rotate API keys every 90 days. Update in Make.com's connection settings without changing the Scenario.

9.2 Data Privacy

Important: When using third-party AI APIs (OpenAI, Anthropic, etc.), email content is transmitted to their servers. Review each provider's data processing agreement (DPA) before processing sensitive business or personal data. Enterprise API plans typically offer stronger data retention controls.

- **Data Minimization** — Send only the email body to the AI; omit attachments and full headers.
- **PII Scrubbing** — Optionally add a Text Parser step to redact phone numbers, SSNs, or credit card patterns before the HTTP call.
- **Retention Policy** — Do not store raw email content in Make.com data stores or Google Sheets beyond what's necessary.
- **Audit Logs** — Make.com logs all executions. On team plans, configure log retention and restrict access to execution history.

9.3 Network Security

HTTPS Enforcement	All Make.com modules use TLS 1.2+ encryption in transit.
IP Allowlisting	If your AI API supports it, allowlist Make.com's IP ranges for additional security.
Webhook Security	If using instant triggers, add a webhook secret header to verify Make.com origin.
GDPR Compliance	Make.com is GDPR-compliant. EU users can select EU data centers in account settings.

10. Testing & Validation

10.1 Development Testing with Make.com

Make.com provides a **Run Once** button that executes the Scenario a single time immediately — bypassing the schedule. This is the primary tool for testing during development. Each module shows a numbered execution bubble with the processed bundle count.

10.2 Test Cases

Test Case	Input Email Type	Expected Outcome	Pass Criteria
TC-01	Short plain text email (< 100 words)	3-line summary generated	Summary delivered in < 10s
TC-02	Long newsletter (> 2000 words)	Truncated, focused summary	No token limit errors
TC-03	HTML-only email (no plain text)	HTML stripped, body summarized	Clean summary, no HTML tags
TC-04	Email with no body (subject only)	Scenario skips gracefully	Filter stops processing
TC-05	Non-English email (Hindi/Spanish)	Summary in same language	AI handles multilingual input
TC-06	AI API returns 429 error	Retry executed, then succeeds	No execution marked failed
TC-07	Gmail OAuth token expires	Break directive fires notification	Alert email sent to admin
TC-08	100 emails arrive simultaneously	All processed as separate bundles	No emails skipped

10.3 Validating AI Summary Quality

Beyond technical testing, evaluate the quality of AI summaries using these criteria:

Quality Metric	How to Measure	Target
Accuracy	Does summary include all key points?	> 95% of key info preserved
Conciseness	Word count of summary vs original	< 20% of original length
Readability	Flesch reading score of summary	Score > 60 (plain English)
Action Clarity	Is the required action clearly stated?	100% of actionable emails flagged
Tone Accuracy	Does tone match the original (urgent/casual)?	Match rate > 90%

11. Performance & Optimization

11.1 Processing Speed

The total processing time per email depends on three factors: Make.com module execution time (~500ms), Gmail API fetch time (~300ms), and AI API inference time (1–8s). Total end-to-end latency is typically **3–10 seconds** per email.

11.2 Operations Budget

Make.com pricing is based on **Operations** — each module execution within a Scenario counts as 1 operation. Plan your operations budget based on email volume:

Make.com Plan	Operations/Month	Est. Emails/Month (6 ops each)	Monthly Cost
Free	1,000	~166 emails	\$0
Core	10,000	~1,666 emails	\$9/mo
Pro	100,000	~16,666 emails	\$16/mo
Teams	500,000	~83,333 emails	\$29/mo

11.3 Optimization Strategies

- **Filter Early** — Add filters immediately after the trigger to skip promotional/spam emails before they consume operations.
- **Batch Processing** — On high-volume inboxes, set 'Max Results' to 10 and process in batches to reduce trigger operations.
- **Caching Labels** — Use Gmail labels to route only specific email categories through the AI summarization path.
- **Minimal HTTP Payload** — Send only the cleaned body text to the AI API; exclude metadata to reduce tokens and cost.
- **Model Tiering** — Use cheap models (Haiku, GPT-4o Mini) for newsletters, premium models (GPT-4o, Sonnet) for client emails.
- **Summary Caching** — Use Make.com Data Store to cache summaries by email hash; avoid re-summarizing forwarded duplicates.
- **Parallel Routes** — Use Make.com's Router module to simultaneously deliver the summary via email AND log it to Sheets.

12. Troubleshooting Guide

■ Scenario Not Triggering

Symptoms	Likely Causes	Fixes
----------	---------------	-------

- No emails processed even after new emails arrive - Scenario shows 0 operations in history	- Gmail connection not authorized - Wrong label filter - Scenario is turned OFF	- Go to Connections → Re-authorize Gmail - Check label name matches exactly (case-sensitive) - Toggle Scenario to ON in the dashboard
--	---	---

■ Empty Email Body in Make.com

Symptoms	Likely Causes	Fixes
- HTTP module receives empty text - Summary is blank or 'No content found'	- Email is image-only - Email uses non-standard multipart MIME encoding	- Add filter: skip if Body(Plain Text) is empty AND Body(HTML) is empty - Use regex in Text Parser to extract body from raw MIME if needed

■ AI API Returns Error

Symptoms	Likely Causes	Fixes
- HTTP module status code 400, 401, 429, 500 - Scenario marked as error/warning	- Invalid API key - Model name typo in JSON body - Request body JSON malformed	- Verify API key in Make.com Data Store - Double-check model name (e.g., 'gpt-4o-mini', not 'gpt4-mini') - Use Make.com's JSON validator in the HTTP module body field

■ Summary Email Not Delivered

Symptoms	Likely Causes	Fixes
- AI summary generated but no email received - Gmail Send module shows success but email missing	- Email sent to wrong address - Gmail spam filter catching automated emails - Missing 'To' field in Gmail Send configuration	- Check 'To' field mapping in Gmail Send module - Whitelist Make.com's sending address in Gmail settings - Check Spam/Sent folder

13. Advanced Enhancements

13.1 Multi-Label Routing with Router Module

Use Make.com's **Router** module to create multiple processing paths based on email properties. For example, route *urgent* emails to a detailed GPT-4o summary, newsletters to a cheaper model, and client emails to a CRM integration:

```
Route 1 Filter: from contains "client.com"      → GPT-4o + CRM Update
Route 2 Filter: subject contains "Newsletter"   → GPT-4o Mini + Archive
Route 3 Filter: subject contains "URGENT"       → GPT-4o + SMS Alert
Route 4 (fallback): All other emails           → GPT-4o Mini + Label Only
```

13.2 Conversation Thread Summarization

Summarize entire email threads (not just individual messages) by using the **Gmail – Get a Thread** module, then iterate over all messages to combine their bodies before sending to the AI:

- Use **Gmail – Search Emails** to find thread by threadId.
- Use **Array Aggregator** to combine all message bodies in the thread.
- Prepend thread context to the AI prompt: *'Summarize this entire email conversation:'*
- Include participant names and dates for richer context summaries.

13.3 Sentiment & Priority Detection

Extend the AI prompt to detect email sentiment and urgency, then auto-tag emails in Gmail based on the AI's assessment:

```
Prompt addition:
"Also respond with these fields in your JSON:",
"sentiment": "positive | negative | neutral",
"urgency": "high | medium | low",
"category": "sales | support | internal | newsletter | spam"
```

Use the **Gmail – Modify Labels** module with conditions based on the JSON output to automatically apply Gmail labels: ■ Urgent, ■ Needs Reply, ■ Info Only.

13.4 Daily Digest Mode

Instead of processing every email individually, configure a daily digest: collect all emails from the day using **Gmail – Search Emails** scheduled at 6 PM, aggregate all bodies with an **Array Aggregator**, and send a single AI-generated daily digest email with a table of all email summaries.

13.5 Integration with Google Sheets Dashboard

Log every summarized email to a Google Sheet for analytics and reporting. Track: sender domain, email volume by category, average response urgency, and AI API cost per day. Build charts inside Sheets to visualize email patterns over time.

14. Use Cases & Business Impact

Executive Email Management	C-suite and managers receive hundreds of emails daily. The agent auto-summarizes all inbound emails before the executive's morning check-in, delivering a digest of action items.	✓ 90-minute daily time saving ✓ Zero missed action items ✓ Prioritized inbox view
Customer Support Triage	Support teams use the agent to pre-summarize customer emails before they reach agents. Sentiment detection routes high-frustration tickets to senior agents automatically.	✓ 40% faster initial response ✓ Reduced agent cognitive load ✓ Automatic escalation routing
Sales Team Intelligence	Sales reps receive summaries of prospect replies instantly, with the AI identifying buying signals, objections, and next steps from email conversations.	✓ Faster lead response times ✓ Consistent follow-up actions ✓ CRM auto-population
Legal Document Monitoring	Legal teams monitor inbound contracts and legal notices. Long agreements are summarized to highlight key clauses, deadlines, and required signatures.	✓ Risk flagging on key terms ✓ Deadline alerts ✓ Review time reduced by 60%
Newsletter Intelligence	Research and marketing teams subscribe to dozens of industry newsletters. The agent distills each one to 3 bullet points, delivered as a single daily digest.	✓ Stay informed without reading ✓ Competitive intelligence ✓ Market trend tracking
HR & Recruitment Inbox	HR teams summarize incoming job applications and candidate emails, extracting qualifications, availability, and salary expectations automatically.	✓ Faster candidate screening ✓ Standardized evaluation format ✓ Reduced admin time

14.1 ROI Analysis

Metric	Before Agent	After Agent	Improvement
Time reading emails/day	2.6 hours	0.5 hours	81% reduction
Missed action items/week	4–7 items	0–1 items	~95% reduction
Email response time	4–8 hours	< 30 minutes	87% faster
AI API cost/month	N/A	\$2–\$15/mo	< \$0.50/1000 emails
Make.com cost/month	N/A	\$9–\$29/mo	Fixed, predictable

15. Conclusion

The **AI Email Summary Agent** built on Make.com represents a powerful, accessible, and cost-effective solution to the universal problem of email overload. By combining Make.com's native Gmail modules with the flexibility of the HTTP module and the intelligence of modern AI language models, this agent delivers real productivity gains without requiring any code, servers, or deep technical expertise.

The architecture described in this document covers every layer of the system: from OAuth authentication and email parsing, through prompt engineering and AI model selection, to error handling, security hardening, and business use case deployment. The modular design of the Make.com Scenario means that each component can be independently modified, upgraded, or replaced as your needs evolve.

15.1 Key Takeaways

- **Make.com** provides a zero-code, fully managed platform ideal for email automation at scale.
- **Gmail Modules** (Watch Emails, Get Email, Send Email) form the core input/output of the agent.
- **HTTP Module** provides universal AI API connectivity — works with any LLM provider.
- **Prompt Engineering** is the primary lever for controlling summary quality and format.
- **Error Handling** (Retry, Ignore, Break) ensures production reliability without manual intervention.
- **Security** is maintained through OAuth 2.0, encrypted key storage, and minimal data transmission.
- **ROI is immediate** — most users recoup the Make.com subscription cost in the first day of use.

15.2 Next Steps

Step 1 — Deploy	Create the Make.com Scenario using the module sequence in Section 3.2.
Step 2 — Connect	Authorize your Gmail account and add your AI API key to Make.com's Data Store.
Step 3 — Test	Run the Scenario with test emails from Section 10.2 and validate output quality.
Step 4 — Tune	Adjust the AI prompt using the templates in Section 6.2 for your email types.
Step 5 — Scale	Enable advanced routing (Section 13.1) and Google Sheets logging (Section 13.5).
Step 6 — Monitor	Review weekly execution history; optimize operations budget as volume grows.

■ **You're ready to build!** This document has covered everything from initial setup to advanced enterprise patterns. Your AI Email Summary Agent on Make.com will be live and saving you hours each week. The future of your inbox is automated, intelligent, and effortless.
